

Machine Learning – Life Cycle Tools

Benjamin Wagner

Gliederung

- **MLOps**
- **ML Prozessmodelle**
 - Crips DM
 - Modell nach Google
- **Komponenten eines vollständigen MLOps Systems**
- **Technologie-Stack**
 - Single Solutions
 - Integrated Solutions



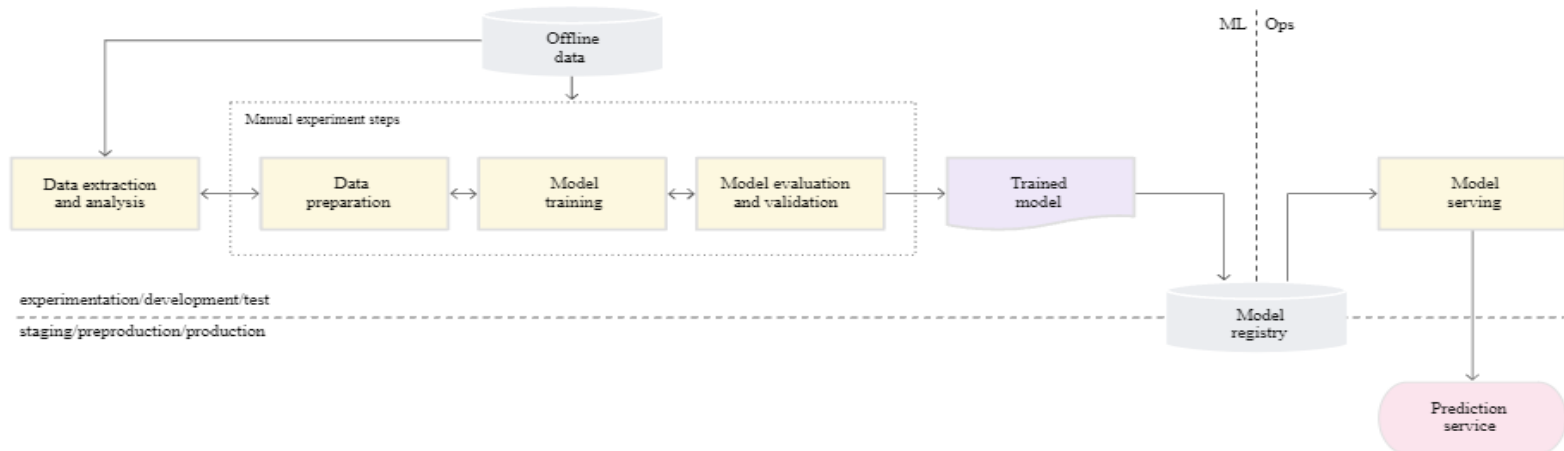

MLOps

Anwendung der DevOps Prinzipien auf Machine Learning Systeme

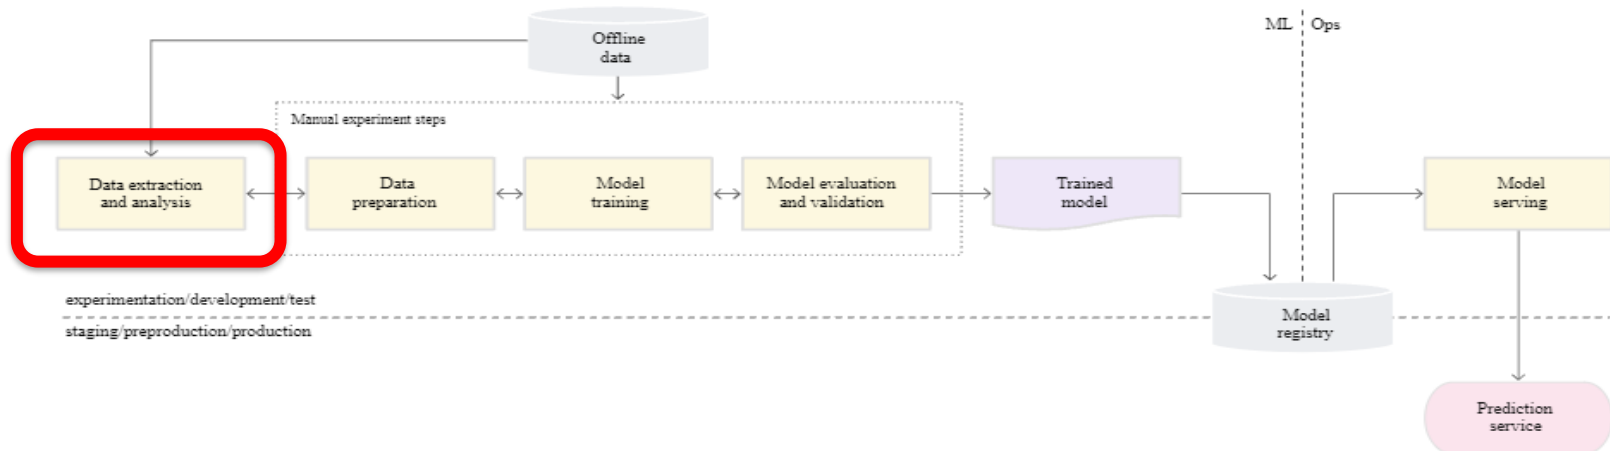

- Reproduzierbarkeit
- Skalierbarkeit
- Automation
- Deployment
- Monitoring und Management
- ...

ML Prozessmodelle

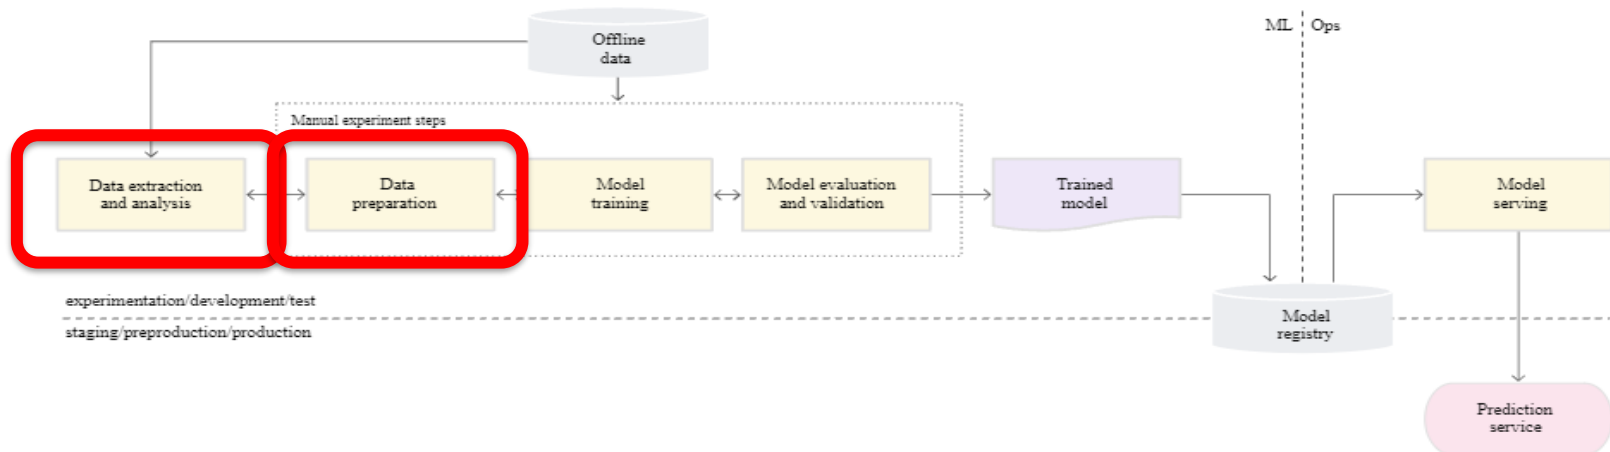

Crisp DM / Google



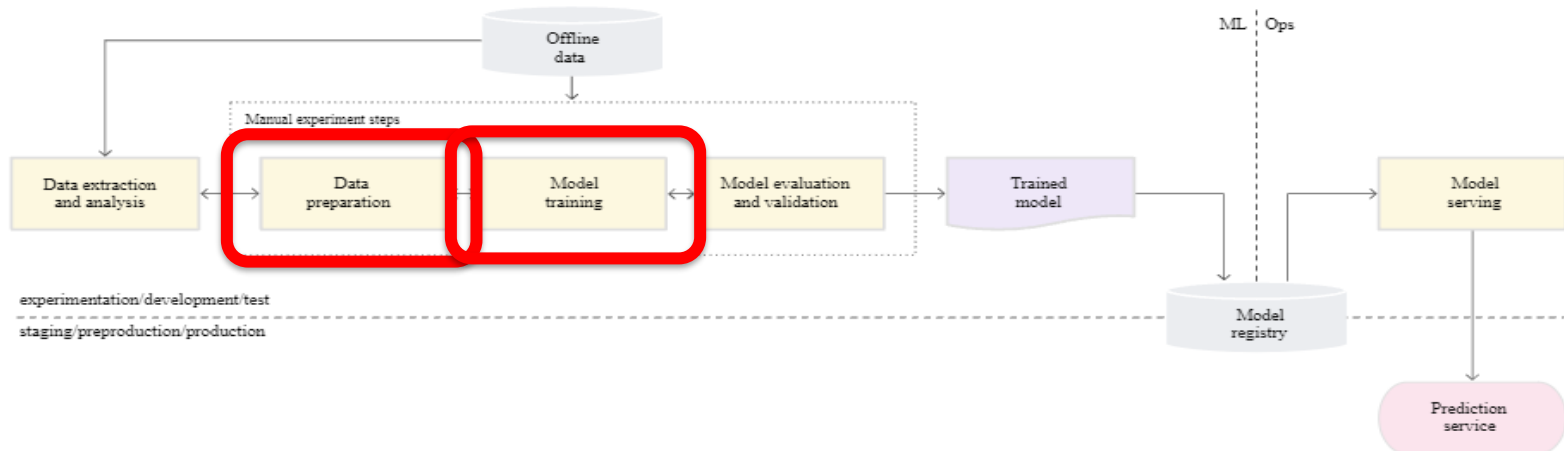

Crisp DM / Google



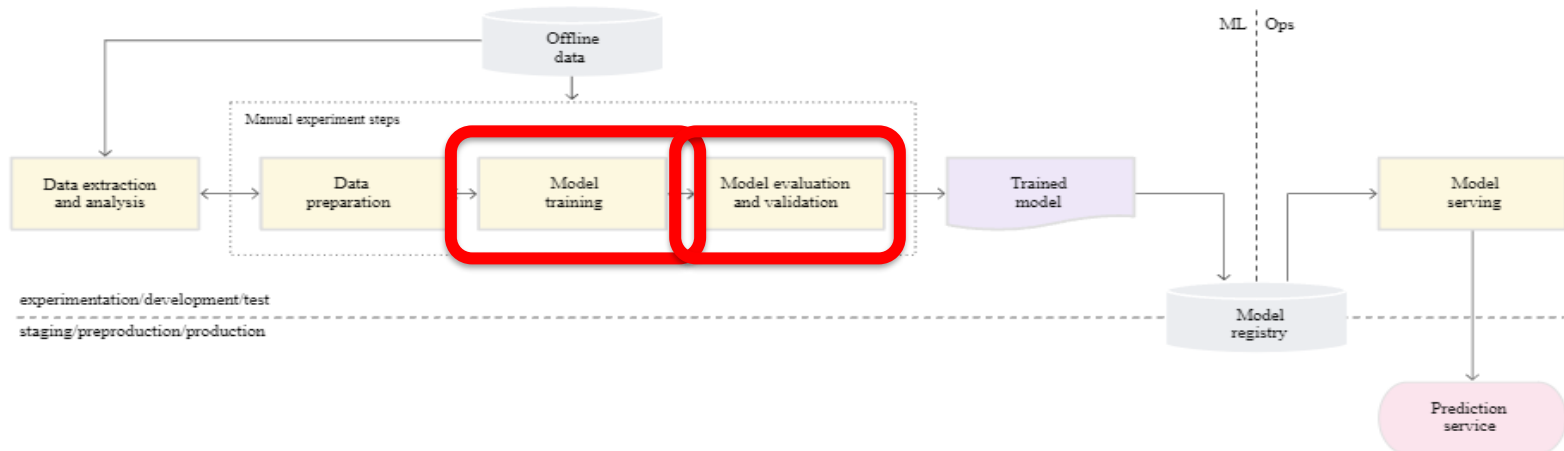

Crisp DM / Google



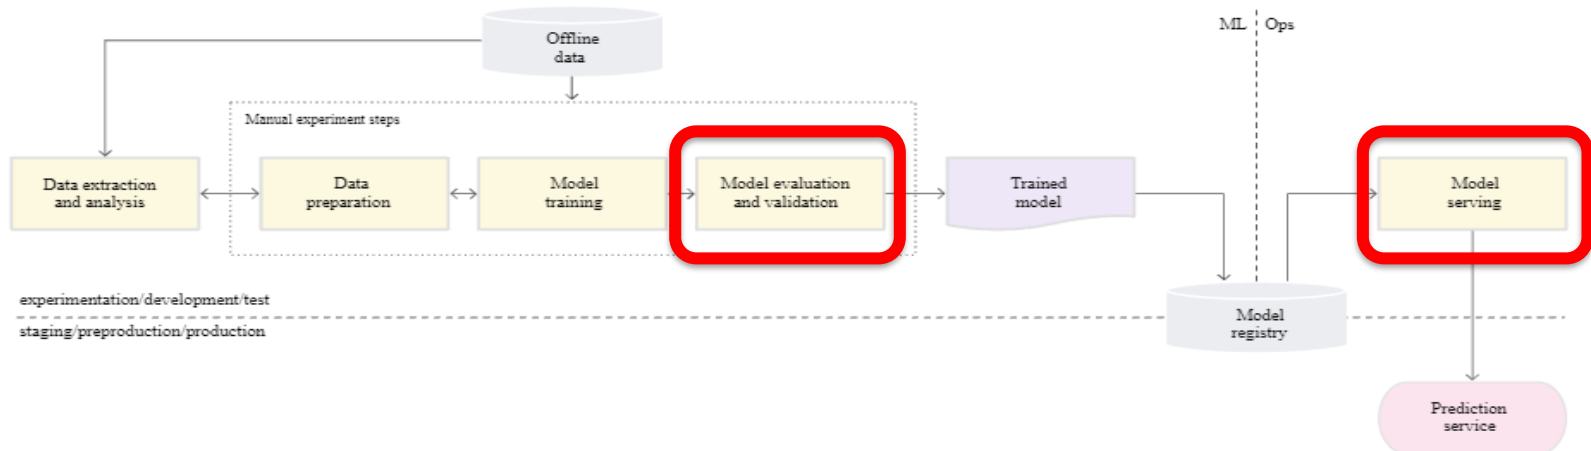

Crisp DM / Google



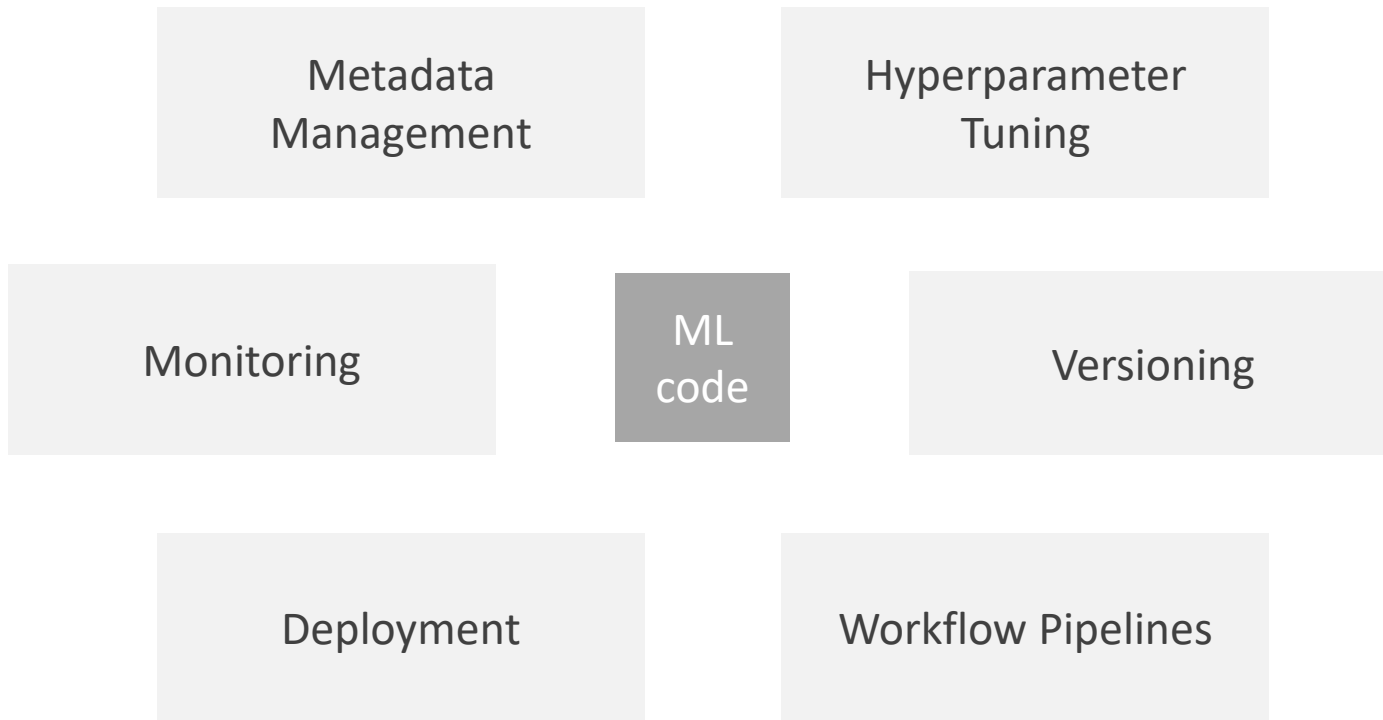
Crisp DM / Google

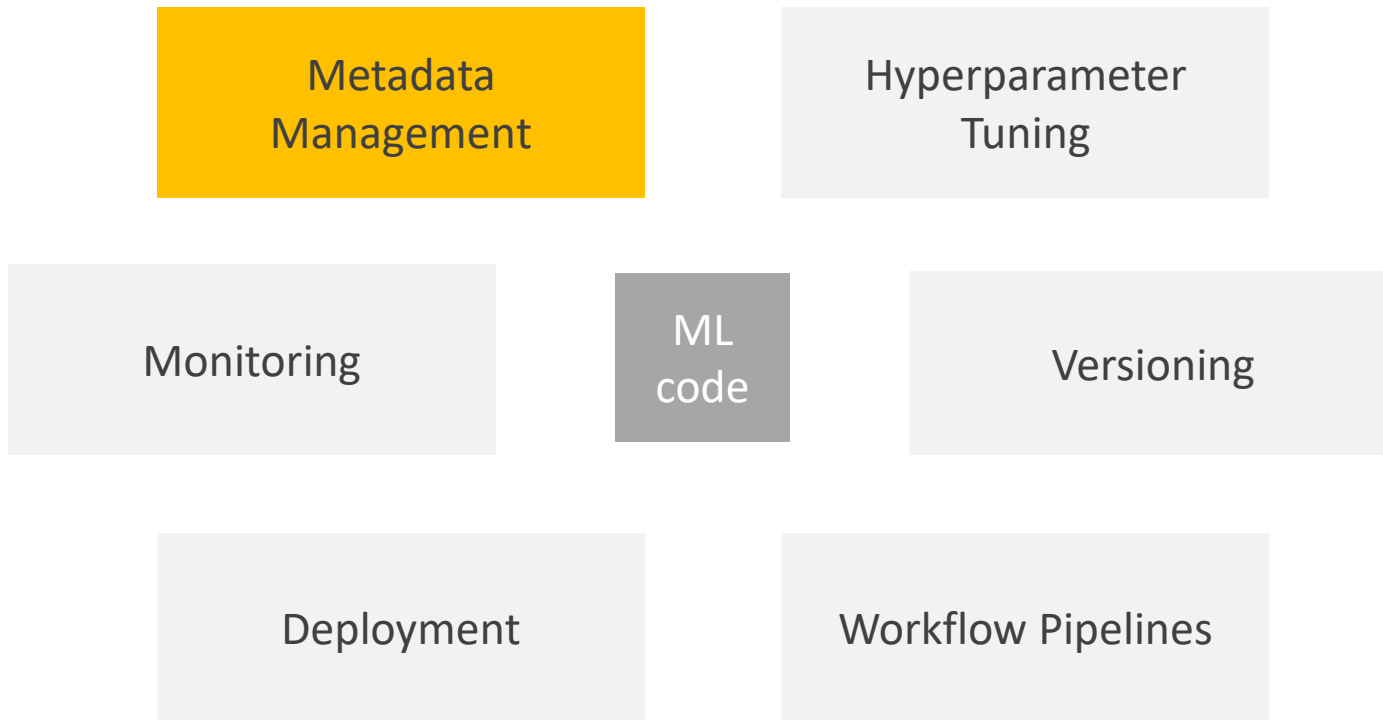


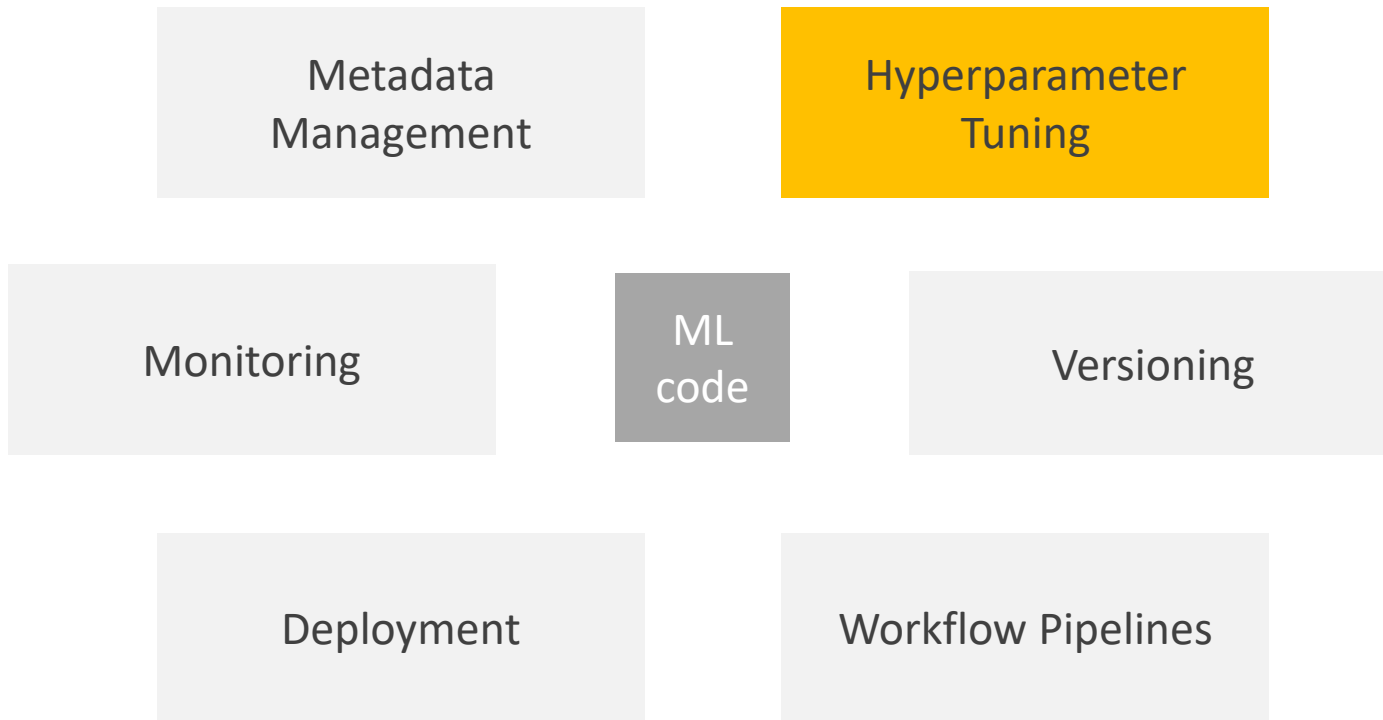
Crisp DM / Google

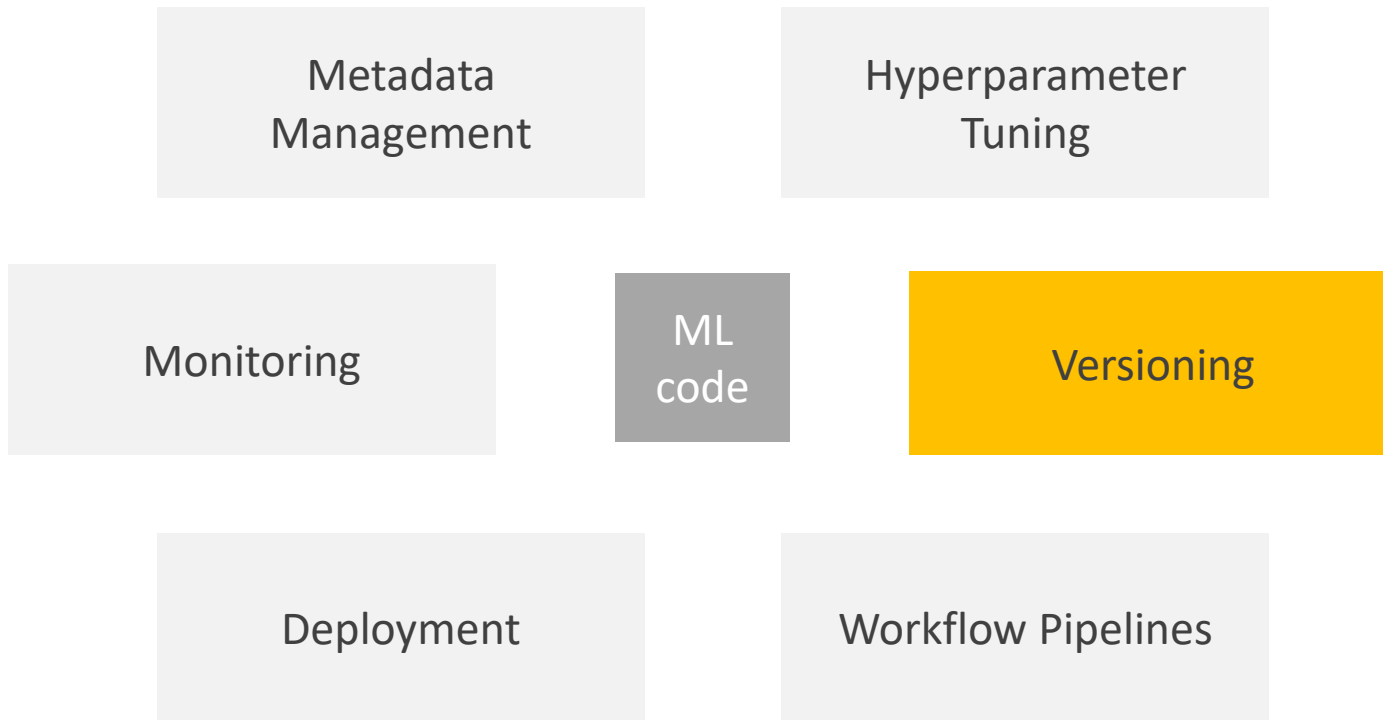


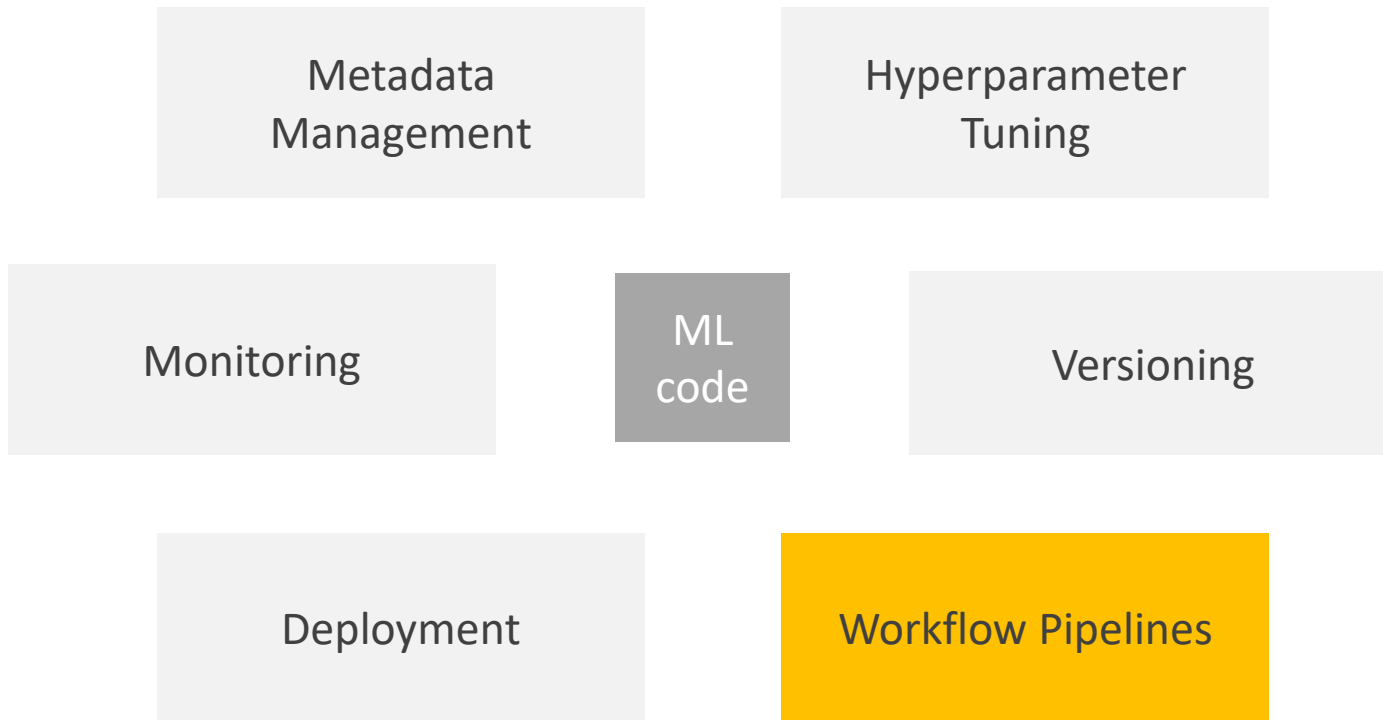
Komponenten eines Vollständigen MLOps System

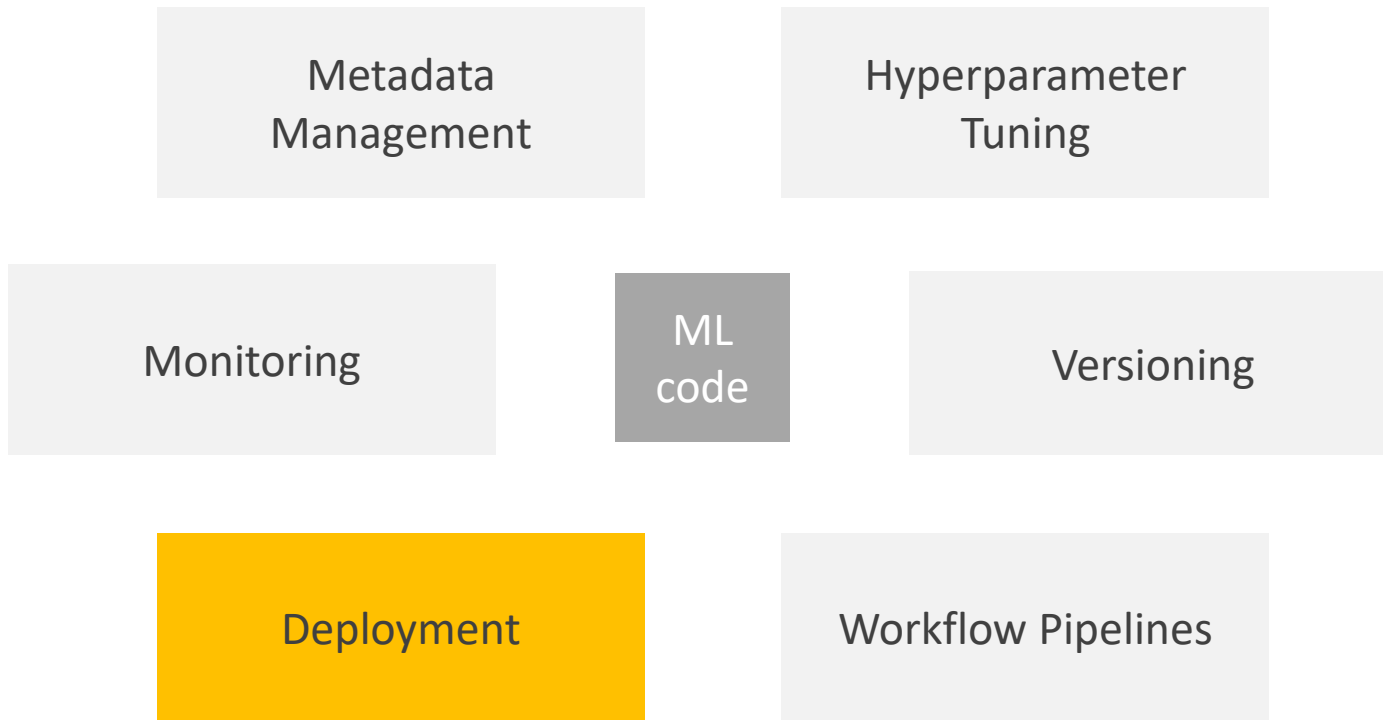


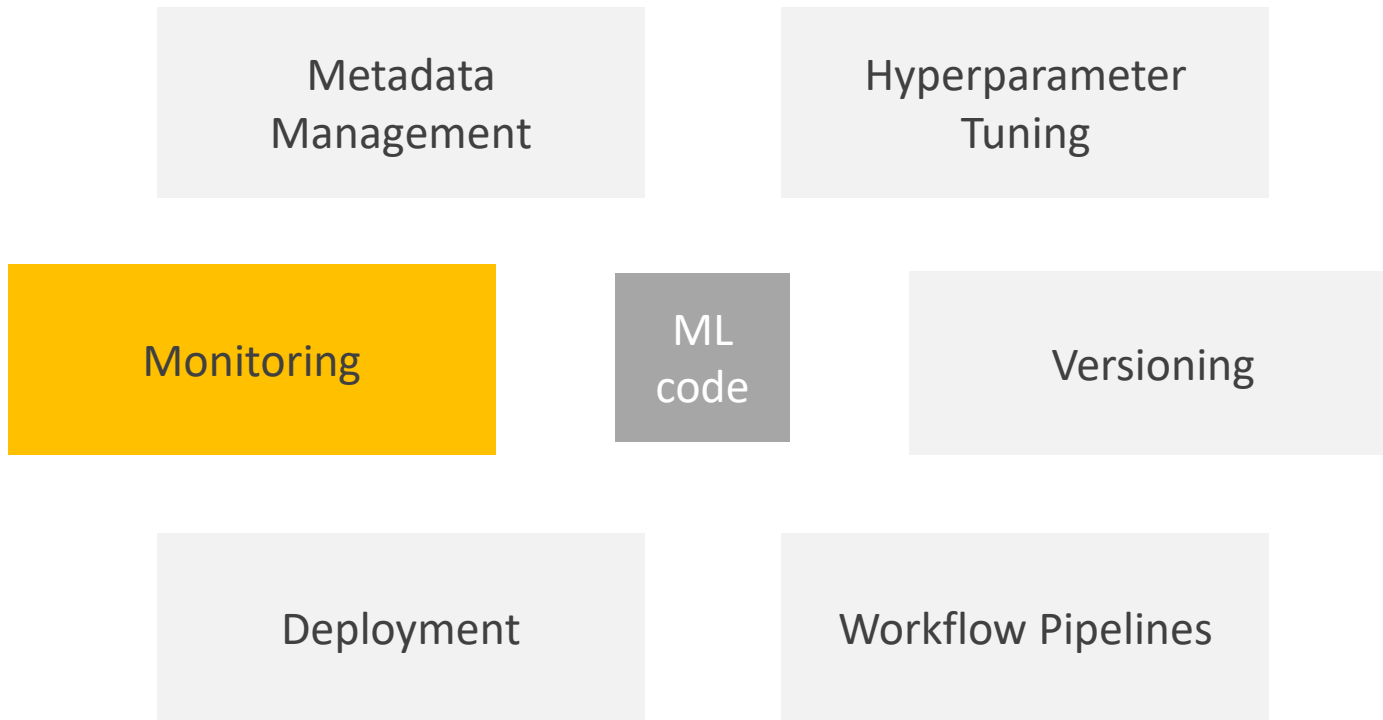












Technologie-Stack

Single Solutions

Metadata management

- ML Flow
- Neptune
- Comet

Hyperparameter Tuning

- Optuna
- SigOPT

Monitoring

- Fiddler
- Amazon Sage Maker

ML code

Versioning

- Pachyderm
- Apache Airflow
- DVC
- GIT

Deployment

- BentoML
- Cortex

Workflow Pipelines

- Kubeflow
- Polyaxon

Metadata Management

ML Flow

ML Flow

Metadata Management

Metadata
Management

Hyperparameter
Tuning

Monitoring

ML
code

Versioning

Deployment

Workflow Pipelines

mlflow

Experiments

Models

GitHub Docs

Experiments

Search Experiments

Default

Default

Track machine learning training runs in an experiment. [Learn more](#)

Experiment ID: 0

Notes

Showing 5 matching runs

[Refresh](#) [Compare](#) [Delete](#) [Download CSV](#) [Start Time](#) [All](#)

[Columns](#) [Only show differences](#) [metrics.rmse < 1 and params.model = "tree"](#) [Search](#) [Filter](#) [Clear](#)

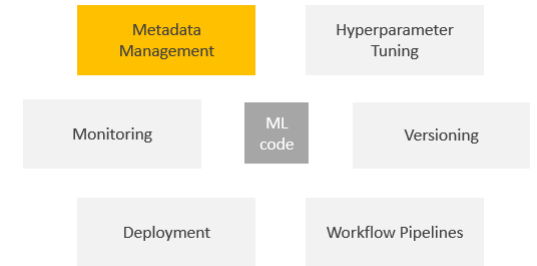
								Metrics		Parameters >		
<input type="checkbox"/>	Start Time	Duration	Run Name	User	Source	Version	Models	accuracy	loss	batch_size	class_weight	epochs
<input type="checkbox"/>	4 minutes ago	424s	-	Benji	main.py	-	keras	0.994	0.017	None	None	10
<input type="checkbox"/>	5 minutes ago	15.7s	-	Benji	main.py	-	keras	0.975	0.078	None	None	3
<input type="checkbox"/>	6 minutes ago	15.2s	-	Benji	main.py	-	keras	0.97	0.106	None	None	3
<input type="checkbox"/>	6 minutes ago	8.3s	-	Benji	main.py	-	keras	0.906	0.339	None	None	1
<input type="checkbox"/>	8 minutes ago	6.6s	-	Benji	main.py	-	keras	0.806	0.771	None	None	1

Load more

ML Flow

Metadata Management

Metrics




Name	Value
accuracy 	0.994
loss 	0.017

ML Flow

Metadata Management

Environment Dependencies

```
channels:  
- conda-forge  
dependencies:  
- python=3.9.5  
- pip  
- pip:  
  - mlflow  
  - keras==2.6.0  
  - pillow==8.4.0  
  - scipy==1.7.1  
  - tensorflow==2.6.0  
name: mlflow-env
```



ML Flow

Metadata Management

Model Summary

Metadata
Management

Hyperparameter
Tuning

Monitoring

ML
code

Versioning

Deployment

Workflow Pipelines

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 128)	100480
dense_1 (Dense)	(None, 128)	16512
dense_2 (Dense)	(None, 10)	1290

Total params: 118,282

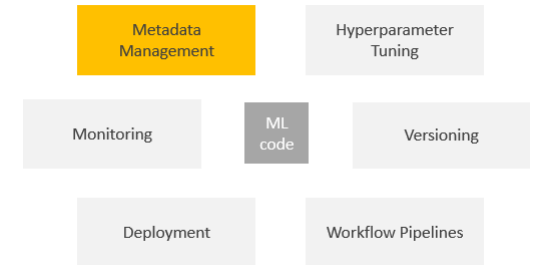
Trainable params: 118,282

Non-trainable params: 0

ML Flow

Metadata Management

- **Läuft mit Python und R**
- **Sehr einfach einzubinden in bestehende Projekte**
 - Z.b via `>> mlflow.tensorflow.autolog()`
- **Unterstützt diverse Bibliotheken**
 - Tensorflow, Scikit, Pytorch, Keras, Fastai ...
- **Übersichtliches UI**
 - Starten `>> mlflow ui`


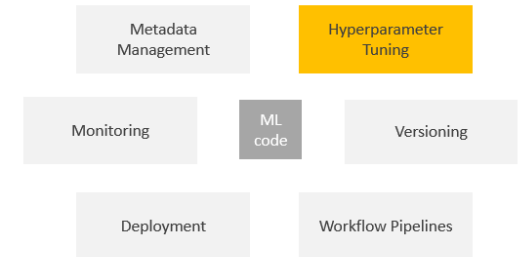


Hyperparameter Tuning

Optuna

Optuna

Hyperparameter Tuning



- 1. Aufstellen des Suchraumes**
- 2. Abbrechen von wenig erfolgsversprechenden Strategien**

Code Integration

```
import optuna
```

```
def objective(trial):
```

Your code here!

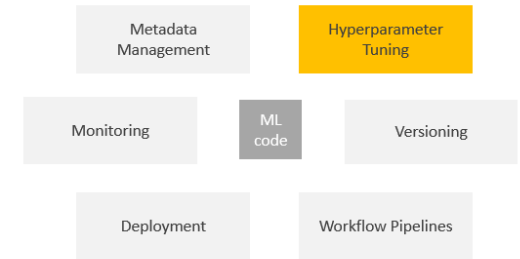
```
    return evaluation_score
```

```
study = optuna.create_study()
```

```
study.optimize(objective, n_trials= Number of trials )
```

Optuna


Hyperparameter Tuning



Code Integration

```
657/657 [=====] - 1s 895us/step - loss: 0.0939 - accuracy: 0.9719
0.9719047546386719
0.09385089576244354
[I 2021-11-03 15:55:32,337] Trial 99 finished with value: 0.09385089576244354 and parameters: {'n_layers': 1, 'n_units_l0': 237, 'optimizer': 'Adam', 'n_epochs': 7}. Best is trial 31 with value: 0.08841179311275482.
Process finished with exit code 0
|
```

Parameter Einfluss



Optuna

Hyperparameter Tuning

Metadata
Management

Hyperparameter
Tuning

Monitoring

ML
code

Versioning

Deployment

Workflow Pipelines

Parallel Computing

```
$ python example.py
[I 2019-05-21 11:16:43,493] Using an existing study with name 'example-study' instead of creating a new one.
[I 2019-05-21 11:16:53,916] Finished trial#9 resulted in value: 0.744140625. Current best value is 0.744140625 with parameters: {'momentum_sgd_lr': 7.944022303405195e-05, 'n_layers': 1, 'n_units_l0': 28.753028463759527, 'optimizer': 'MomentumSGD', 'weight_decay': 2.1924271434430527e-06}.
total [#####] 58.88%
this epoch [#####] 88.80%
138 iter, 5 epoch / 10 epochs
17.446 iters/sec. Estimated time to finish: 0:00:05.524267.
```

```
$ python example.py
[I 2019-05-21 11:16:43,621] Using an existing study with name 'example-study' instead of creating a new one.
[I 2019-05-21 11:16:55,890] Finished trial#10 resulted in value: 0.8976862980052829. Current best value is 0.16624098271131516 with parameters: {'momentum_sgd_lr': 0.012378325131946236, 'n_layers': 3, 'n_units_l0': 10.118405976356666, 'n_units_l1': 99.22366485593328, 'n_units_l2': 6.970853785986278, 'optimizer': 'MomentumSGD', 'weight_decay': 0.0007099850265678694}.
total [#####] 81.49%
this epoch [#####] 14.93%
191 iter, 8 epoch / 10 epochs
25.217 iters/sec. Estimated time to finish: 0:00:01.720039.
```

```
$ python example.py
[I 2019-05-21 11:16:42,729] Using an existing study with name 'example-study' instead of creating a new one.
[I 2019-05-21 11:16:54,781] Finished trial#8 resulted in value: 0.9162409855052829. Current best value is 0.16624098271131516 with parameters: {'momentum_sgd_lr': 0.012378325131946236, 'n_layers': 3, 'n_units_l0': 10.118405976356666, 'n_units_l1': 99.22366485593328, 'n_units_l2': 6.970853785986278, 'optimizer': 'MomentumSGD', 'weight_decay': 0.0007099850265678694}.
```


```
instead of creating a new one.
[I 2019-05-21 11:16:51,250] Finished trial#7 resulted in value: 0.80581430737. Current best value is 0.805814303457737 with parameters: {'adam_alpha': 0.5501424552926803e-05, 'n_layers': 1, 'n_units_l0': 4.237994876174356, 'optimizer': 'Adam', 'weight_decay': 4.2909422937320695e-09}.
[I 2019-05-21 11:17:02,585] Finished trial#11 resulted in value: 0.1422776452606. Current best value is 0.14227764308452606 with parameters: {'momentum_sgd_lr': 0.02227271795421007, 'n_layers': 3, 'n_units_l0': 6.0048260623270, 'n_units_l1': 27.182953480840045, 'n_units_l2': 24.852836817988184, 'optimizer': 'MomentumSGD', 'weight_decay': 1.378459092583759e-10}.
total [#####] 9.81%
this epoch [#####] 98.13%
23 iter, 0 epoch / 10 epochs
26.735 iters/sec. Estimated time to finish: 0:00:07.906416.
```

```
$ python example.py
[I 2019-05-21 11:16:42,230] Using an existing study with name 'example-study' instead of creating a new one.
[I 2019-05-21 11:16:56,449] Finished trial#6 resulted in value: 0.882737371315. Current best value is 0.16624098271131516 with parameters: {'momentum_sgd_lr': 0.012378325131946236, 'n_layers': 3, 'n_units_l0': 10.118405976356666, 'n_units_l1': 99.22366485593328, 'n_units_l2': 6.970853785986278, 'optimizer': 'MomentumSGD', 'weight_decay': 0.0007099850265678694}.
total [#####] 78.93%
this epoch [#####] 89.33%
185 iter, 7 epoch / 10 epochs
24.806 iters/sec. Estimated time to finish: 0:00:01.990426.
```

```
$ python example.py
[I 2019-05-21 11:16:42,022] Using an existing study with name 'example-study' instead of creating a new one.
[I 2019-05-21 11:16:54,576] Finished trial#5 resulted in value: 0.1662409831516. Current best value is 0.16624098271131516 with parameters: {'momentum_sgd_lr': 0.012378325131946236, 'n_layers': 3, 'n_units_l0': 10.118405976356666, 'n_units_l1': 99.22366485593328, 'n_units_l2': 6.970853785986278, 'optimizer': 'MomentumSGD', 'weight_decay': 0.0007099850265678694}.
total [#####] 76.37%
this epoch [#####] 63.73%
179 iter, 7 epoch / 10 epochs
22.063 iters/sec. Estimated time to finish: 0:00:02.509897.
```


Optuna

Hyperparameter Tuning



- Integrierte Bibliotheken

- Pytorch lightning
- Pytorch ignite
- FastAi

- Parallel Computing

- Visualisierung des Suchraumes

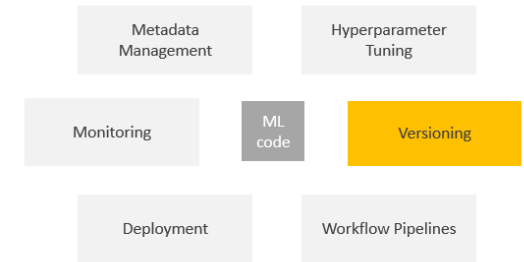
- Gewichtung des Parametereinflusses

Versioning

DVC/MLFlow/Git

Data Version Control

Versioning

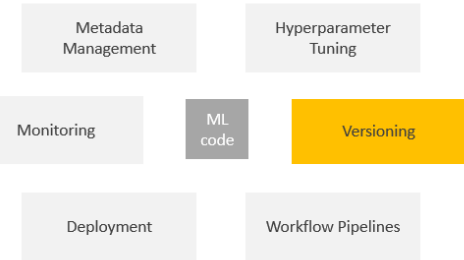


- **Versionierung von Datensätzen**
- **Prinzip: Versionierung des Codes über Git**
- **Datensätze zu groß für GitHub Repos**
- **Trainingsdaten werden über DVC auf einen Storage geladen (z.b. direkt auf gdrive)**
- **Im Git Repository wird eine dvc/config file abgelegt welche das Git-Repo mit dem DVC-Repo verknüpft**

ML Flow

Versioning

Versionierung von Modellen



mlflow Experiments Models GitHub Docs

Experiments + ←

Search Experiments

Default 🔗 📄

i Track machine learning training runs in an experiment. [Learn more](#) ×

Experiment ID: 0

► Notes [🔗](#)

Showing 5 matching runs

↻ Refresh ⚖ Compare 🗑 Delete 📄 Download CSV ⌵ Start Time All

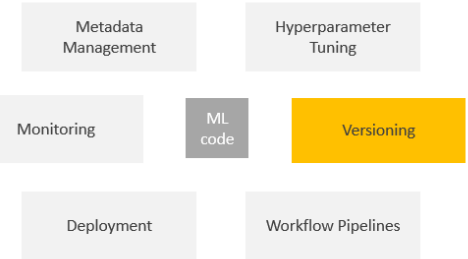
☰ 📄 ⚙ Columns Only show differences 🔍 metrics.rmse < 1 and params.model = "tree" Search ⚙ Filter Clear

								Metrics		Parameters >		
<input type="checkbox"/>	↓ Start Time	Duration	Run Name	User	Source	Version	Models	accuracy	loss	batch_size	class_weight	epochs
<input type="checkbox"/>	🟢 4 minutes ago	42.4s	-	Benji	main.py	-	keras	0.994	0.017	None	None	10
<input type="checkbox"/>	🟢 5 minutes ago	15.7s	-	Benji	main.py	-	keras	0.975	0.078	None	None	3
<input type="checkbox"/>	🟢 6 minutes ago	15.2s	-	Benji	main.py	-	keras	0.97	0.106	None	None	3
<input type="checkbox"/>	🟢 6 minutes ago	8.3s	-	Benji	main.py	-	keras	0.906	0.339	None	None	1
<input type="checkbox"/>	🟢 8 minutes ago	6.6s	-	Benji	main.py	-	keras	0.806	0.771	None	None	1


Load more

GIT

Versioning



Versionierung von Code und Pipelines




Workflow Pipelines

Kedro

Kedro

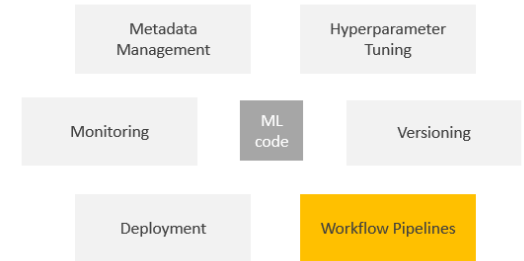
Workflow Pipelines

- **Python library**
- **Erstellen von Node**
 - Funktion
 - Inputs
 - Outputs
 - Name
- **Verbinden von Nodes zu einer Pipeline**
- **Verbinden von mehreren Pipelines**
- **Speichern/Teilen von Pipelines**
- **Visualisieren von Pipelines**



Kedro

Workflow Pipelines



- **Neues Projekt: >> kedro new**
- **Datensets importieren**
- **Catalog definieren**
- **Funktionen/Nodes definieren**
- **Pipelines registrieren**
- **Sequentiell/Parallel laufen lassen**

Kedro

Workflow Pipelines

Datensets/Artefakte definieren

```
companies:
  type: pandas.CSVDataSet
  filepath: data/01_raw/companies.csv
  layer: raw

reviews:
  type: pandas.CSVDataSet
  filepath: data/01_raw/reviews.csv
  layer: raw

shuttles:
  type: pandas.ExcelDataSet
  filepath: data/01_raw/shuttles.xlsx
  layer: raw
```

Metadata
Management

Hyperparameter
Tuning

Monitoring

ML
code

Versioning

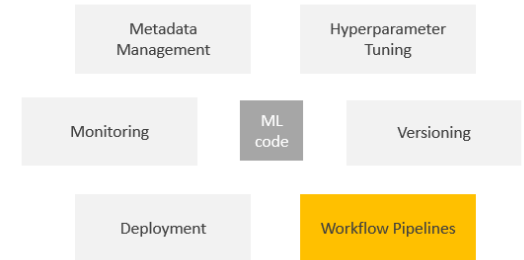
Deployment

Workflow Pipelines

Kedro

Workflow Pipelines

(Node)-Funktion



```
def preprocess_companies(companies: pd.DataFrame) -> pd.DataFrame:
    """Preprocesses the data for companies.

    Args:
        companies: Raw data.
    Returns:
        Preprocessed data, with `company_rating` converted to a float and
        `iata_approved` converted to boolean.
    """
    companies["iata_approved"] = _is_true(companies["iata_approved"])
    companies["company_rating"] = _parse_percentage(companies["company_rating"])
    return companies
```

Kedro

Workflow Pipelines

Metadata
Management

Hyperparameter
Tuning

Monitoring

ML
code

Versioning

Deployment

Workflow Pipelines

Pipeline Definition

```
def create_pipeline(**kwargs):  
    return Pipeline(  
        [  
            node(  
                func=preprocess_companies,  
                inputs="companies",  
                outputs="preprocessed_companies",  
                name="preprocess_companies_node",  
            ),  
            node(  
                func=preprocess_shuttles,  
                inputs="shuttles",  
                outputs="preprocessed_shuttles",  
                name="preprocess_shuttles_node",  
            ),  
            node(  
                func=create_model_input_table,  
                inputs=["preprocessed_shuttles", "preprocessed_companies", "reviews"],  
                outputs="model_input_table",  
                name="create_model_input_table_node",  
            ),  
        ]  
    )
```

Pipelines verbinden

```
def register_pipelines() -> Dict[str, Pipeline]:  
    """Register the project's pipeline.  
  
    Returns:  
        A mapping from a pipeline name to a ``Pipeline`` object.  
  
    """  
    data_processing_pipeline = dp.create_pipeline()  
    data_science_pipeline = ds.create_pipeline()  
  
    return {  
        "__default__": data_processing_pipeline + data_science_pipeline,  
        "dp": data_processing_pipeline,  
        "ds": data_science_pipeline,  
    }
```

Key
Work

Visualisierung

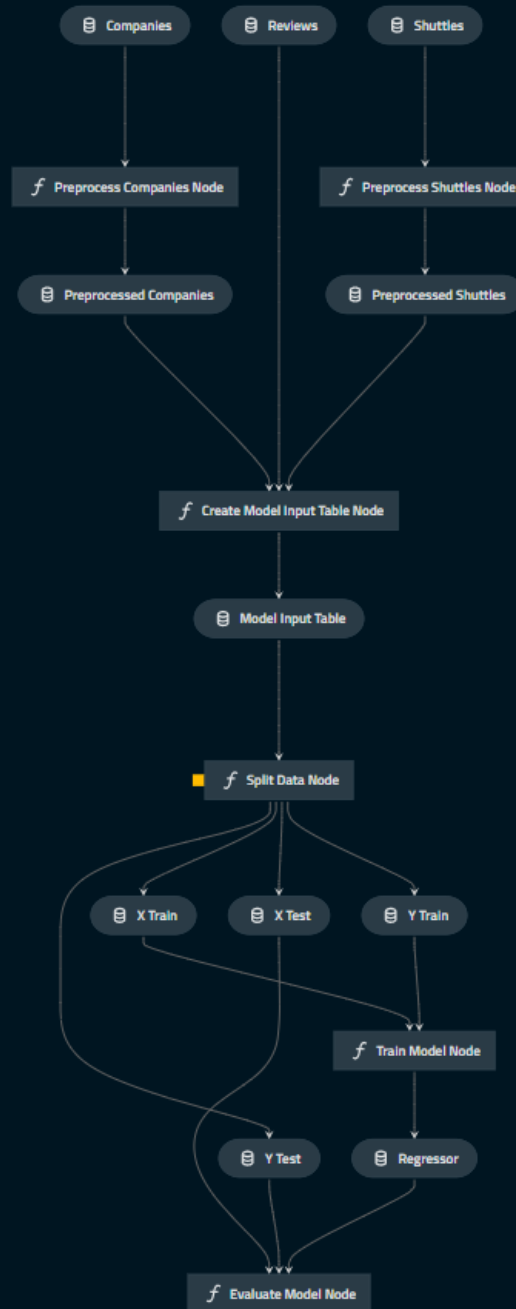
Hochschule

raw

intermediate

primary

models



Hyperparameter
Tuning

IL
de

Versioning

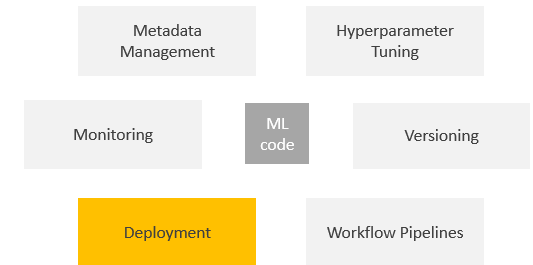
Workflow Pipelines

Deployment

Bento

Bento

Deployment



1. Model

```
# train.py
from sklearn import svm
from sklearn import datasets

# Load training data
iris = datasets.load_iris()
X, y = iris.data, iris.target

# Model Training
clf = svm.SVC(gamma='scale')
clf.fit(X, y)
```

Bento

Deployment

Metadata
Management

Hyperparameter
Tuning

Monitoring

ML
code

Versioning

Deployment

Workflow Pipelines


2. Create API

```
# bento_service.py
import pandas as pd

from bentoml import env, artifacts, api, BentoService
from bentoml.adapters import DataframeInput
from bentoml.frameworks.sklearn import SklearnModelArtifact

@env(infer_pip_packages=True)
@artifacts([SklearnModelArtifact('model')])
class IrisClassifier(BentoService):
    """
    A minimum prediction service exposing a Scikit-Learn model
    """

    @api(input=DataframeInput(), batch=True)
    def predict(self, df: pd.DataFrame):
        """
        An inference API named `predict` with Dataframe input adapter, which codifies
        how HTTP requests or CSV files are converted to a pandas Dataframe object as the
        inference API function input
        """
        return self.artifacts.model.predict(df)
```

3. Pack Model and API

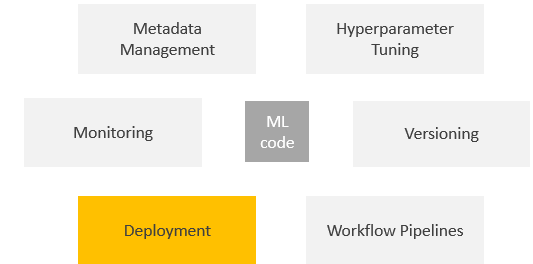
```
# bento_packer.py

# import the IrisClassifier class defined above
from bento_service import IrisClassifier

# Create a iris classifier service instance
iris_classifier_service = IrisClassifier()

# Pack the newly trained model artifact
iris_classifier_service.pack('model', clf)

# Save the prediction service to disk for model serving
saved_path = iris_classifier_service.save()
```



4. Serve API

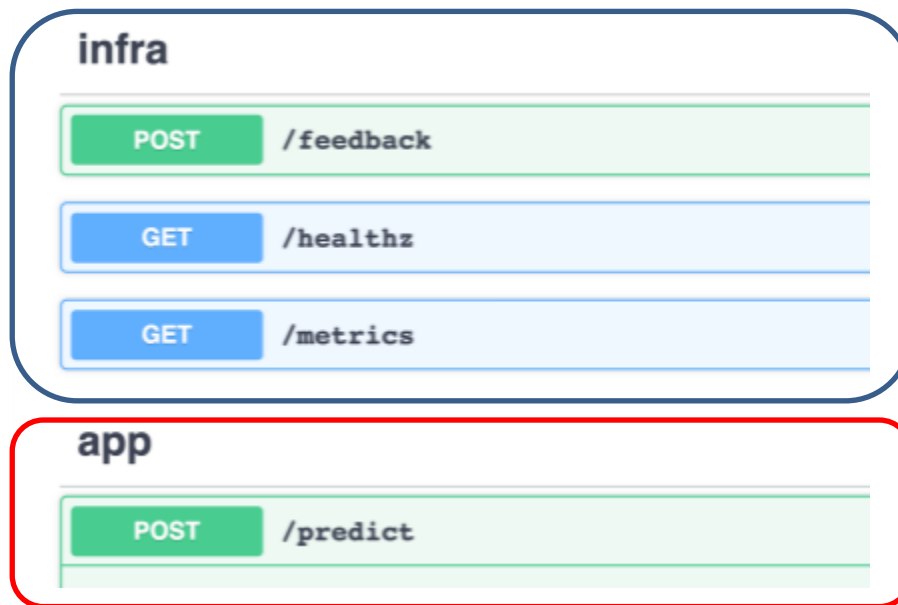
```
bentoml serve IrisClassifier:latest
```

- Model API wird auf localhost:5000 bereitgestellt

Bento

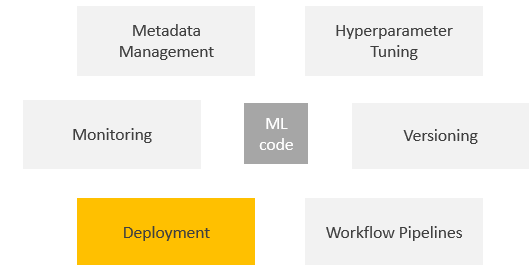
Deployment

Running API Services



API Monitoring

Model Predictions



Bento

Deployment

Metadata
Management

Hyperparameter
Tuning

Monitoring

ML
code

Versioning

Deployment

Workflow Pipelines

Request Body

```
[  
  [5.1, 3.5, 1.4, 0.2],  
  [5.0, 3.5, 1.6, 0.3],  
  [4.9, 4.5, 1.2, 0.2]  
]
```

Execute

Clear

Responses

Curl

```
curl -X POST "http://127.0.0.1:5000/predict" -H "accept: */*" -H "Content-Type: application/json" -d "[  
  [5.1,3.5,1.4,0.2],[5,3.5,1.6,0.3],[4.9,4.5,1.2,0.2]]"
```

Request URL

```
http://127.0.0.1:5000/predict
```

Bento

Deployment

Metadata
Management

Hyperparameter
Tuning

Monitoring

ML
code

Versioning

Deployment

Workflow Pipelines

Response

Code	Details	
200	<div><div>Response body</div><div><pre>[0, 0, 0]</pre><div>Download</div></div><div><div>Response headers</div><div><pre>content-length: 9 content-type: application/json date: Wed, 06 May 2020 00:08:59 GMT request_id: cc0257bd-5299-439e-9f38-27cc731f887c server: Werkzeug/0.16.0 Python/3.7.5</pre></div></div></div>	
Responses		
Code	Description	Links
200	SUCCESS	No links

Monitoring

Fiddler

Fiddler

Monitoring

Model Details

Metadata
Management

Hyperparameter
Tuning

Monitoring

ML
code

Versioning

Deployment

Workflow Pipelines

Properties

BASIC

Fiddler ID

random_forest_26

Name

Random Forest Model

Description

This is models customer bank churn

TECHNICAL

Algorithm

-

Framework

-

Model Task

binary_classification

SCHEMA DETAILS

Input Type

structured

Class Labels

-

Targets

churn(category): [no,yes]

Dataset

[bank_churn](#)

Outputs

1 output(s) found in model schema

NAME ⇅

DATA TYPE ⇅

probability_yes

float

Fiddler

Monitoring

Metadata
Management

Hyperparameter
Tuning

Monitoring


ML
code

Versioning

Deployment


Workflow Pipelines

Model Decisions



Model Accuracy

Accuracy ⓘ



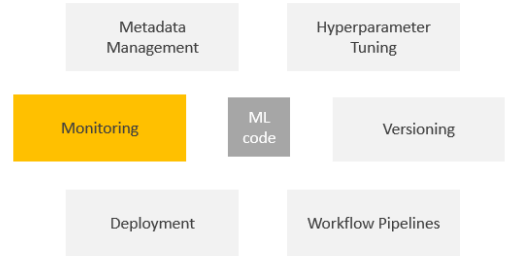
Model Drift



Fiddler

Monitoring

Feature Drift 1. April




Drift Analytics ⓘ

Apr 01, 2021 5PM - 1 Day

Q S

FEATURE ⇅	PREDICTION DRIFT IMPACT ▲ ⓘ	FEATURE DRIFT ⇅ ⓘ	FEATURE IMPACT ⇅ ⓘ
▶ age	45.472%	0.08	33.91%
▶ balance	19.467%	0.09	13.56%
▶ numofproducts	15.149%	0.03	27.13%
▶ isactivemember	9.280%	0.04	14.45%

Feature Drift 12. April




Drift Analytics ⓘ

Apr 12, 2021 5PM - 1 Day

Search

FEATURE ⌵	PREDICTION DRIFT IMPACT ⌵ ⓘ	FEATURE DRIFT ⌵ ⓘ	FEATURE IMPACT ⌵ ⓘ
▶ numofproducts	69.306%	0.32	27.13%
▶ age	15.815%	0.06	33.91%
▶ balance	7.564%	0.07	13.56%
▶ isactivemember	2.299%	0.02	14.45%

Feature Distribution



Technologie-Stack

Integrated Solutions

Azure Machine Learning

Metadata Management

- Python SDK für Azure Machine Learning

Hyperparameter Tuning

- hyperdrive Paket

Monitoring

- Azure Monitor

ML
code

Versioning

- Datasets

Deployment

- Azure Pipelines

Workflow Pipelines

- Azure Machine Learning Pipelines
- Azure Data Factory-Pipelines
- Azure Pipelines

Amazon SageMaker

Prepare →

SageMaker Ground Truth
Label training data for machine learning

SageMaker Data Wrangler NEW
Aggregate and prepare data for machine learning

SageMaker Processing
Built-in Python, BYO R/Spark

SageMaker Feature Store NEW
Store, update, retrieve, and share features

SageMaker Clarify NEW
Detect bias and understand model predictions

Build →

SageMaker Studio Notebooks
Jupyter notebooks with elastic compute and sharing

Built-in and Bring-your-own Algorithms
Dozens of optimized algorithms or bring your own

Local Mode
Test and prototype on your local machine

SageMaker Autopilot
Automatically create machine learning models with full visibility

SageMaker JumpStart NEW
Pre-built solutions for common use cases

Train & tune →

One-click Training
Distributed infrastructure management

SageMaker Experiments
Capture, organize, and compare every step

Automatic Model Tuning
Hyperparameter optimization

Distributed Training Libraries NEW
Training for large datasets and models

SageMaker Debugger NEW
Debug and profile training runs

Managed Spot Training
Reduce training cost by 90%

Deploy & manage →

One-click Deployment
Fully managed, ultra low latency, high throughput

Kubernetes & Kubeflow Integration
Simplify Kubernetes-based machine learning

Multi-Model Endpoints
Reduce cost by hosting multiple models per instance

SageMaker Model Monitor
Maintain accuracy of deployed models

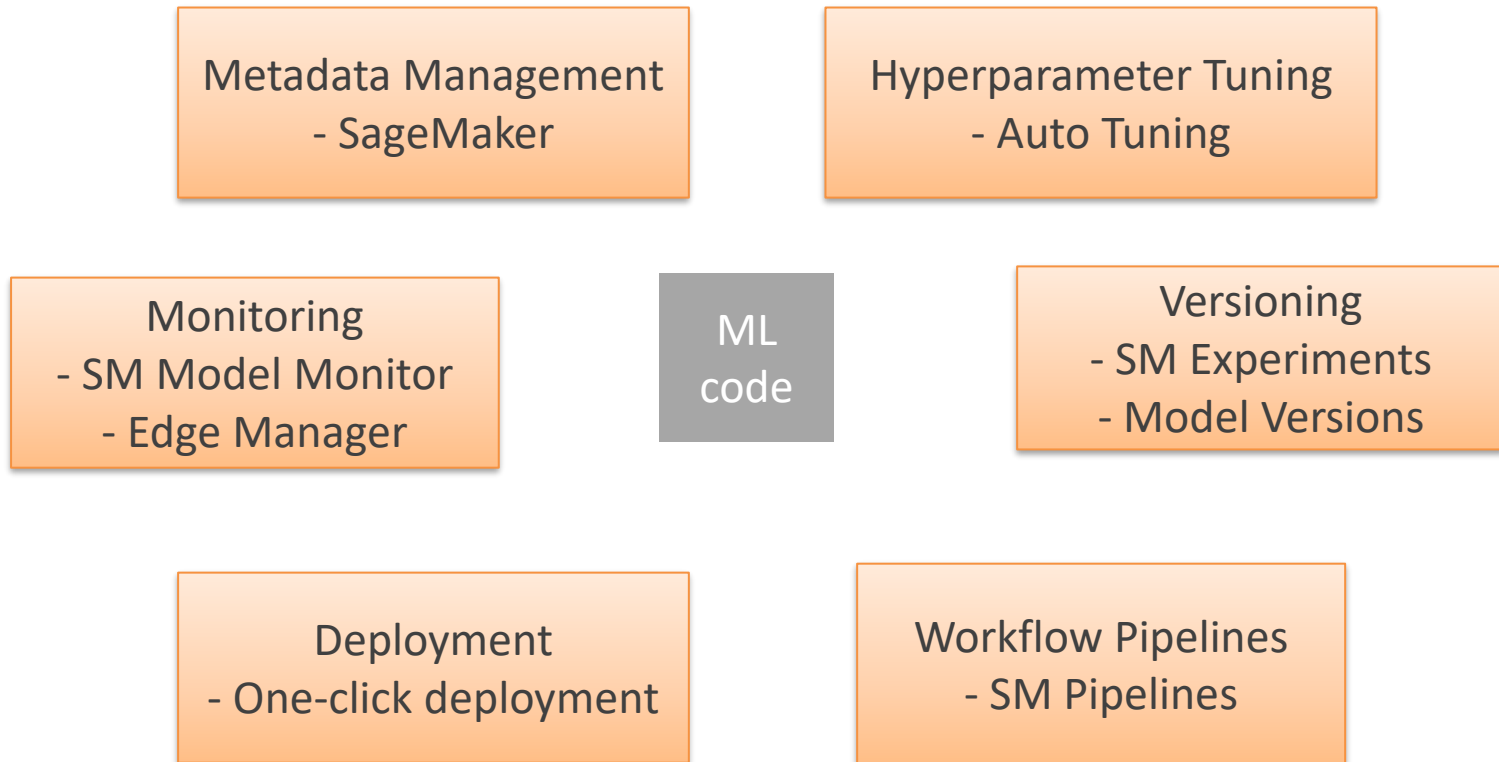
SageMaker Edge Manager NEW
Manage and monitor models on edge devices

SageMaker Pipelines NEW
Workflow orchestration and automation

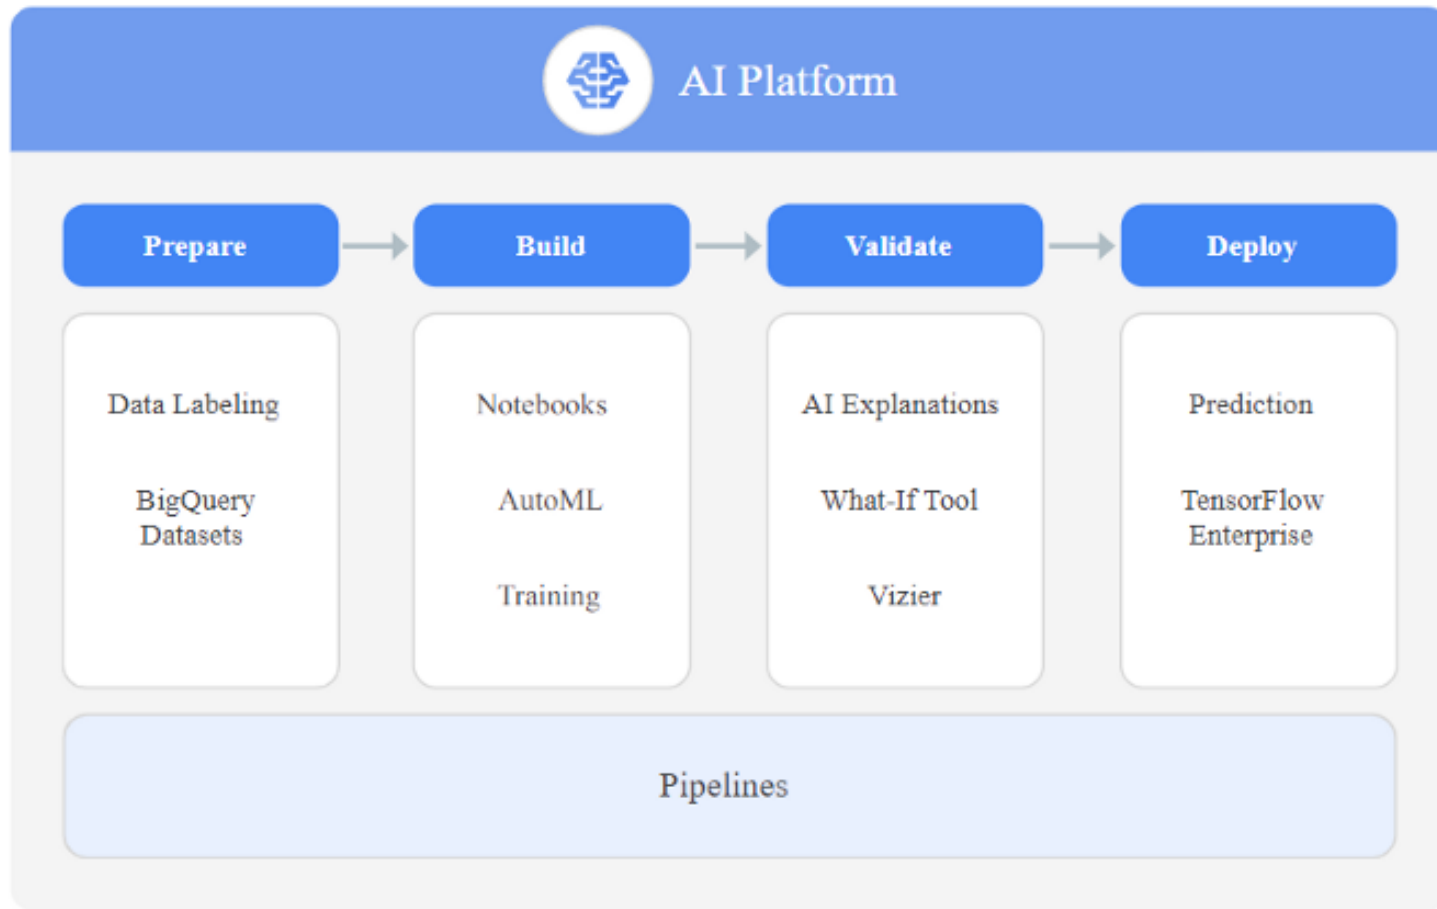
SageMaker Studio

Integrated development environment (IDE) for ML

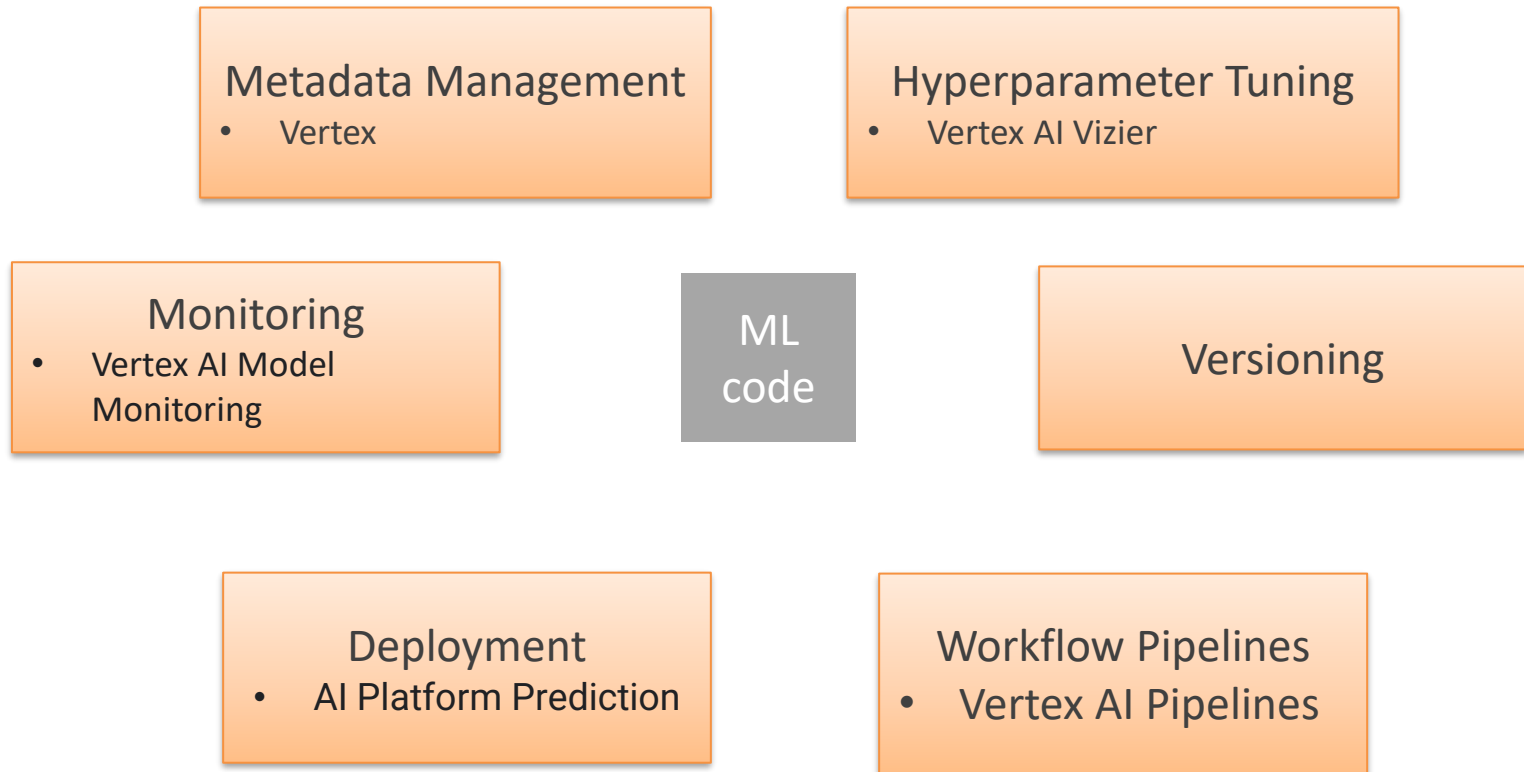
Amazon SageMaker – IDE for ML



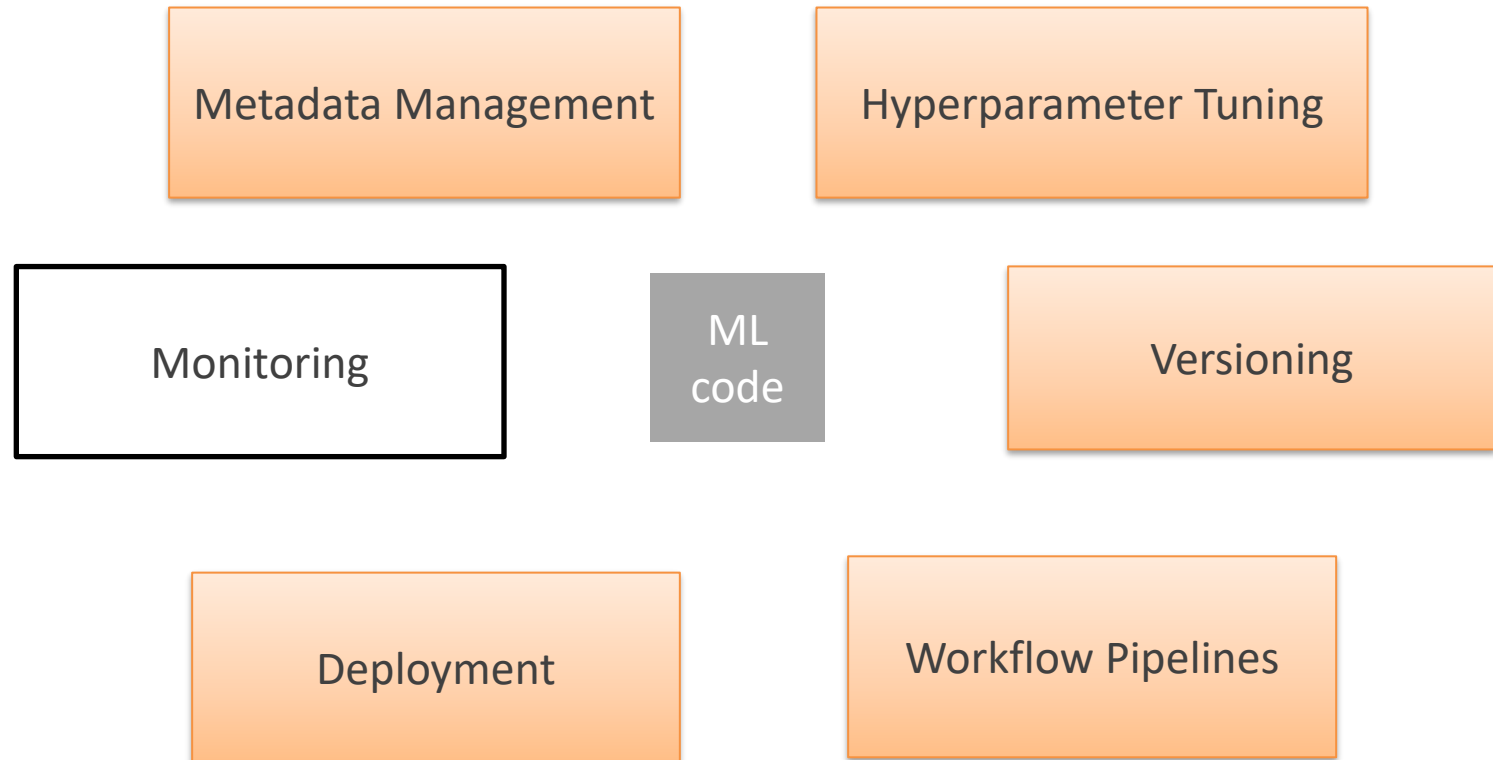
Google Cloud AI



Google Cloud AI - Vertex





TensorFlow Extended



Handlungsempfehlung

- **Flexibilität**
 - Neues Gebiet > Wechsel eines Tools muss einfach möglich sein
- **Stabilität**
 - Untergang von kleinen Tools auf lange Sicht möglich > Auf große Player setzen
- **Priorisierung**
 - Komplexes Gebiet > Da anfangen wo der größte Nutzen liegt





Quellenverzeichnis

- <https://www.trustradius.com/mlops>
- <https://neptune.ai/blog/best-mlops-tools>
- https://cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning#mlops_level_0_manual_process
- <https://optuna.readthedocs.io/en/stable/tutorial/index.html>
- <https://www.youtube.com/watch?v=kLKBcPonMYw>
- <https://towardsdatascience.com/google-cloud-ai-platform-hyper-accessible-ai-machine-learning-cddd8c3348b3>
- <https://searchcloudcomputing.techtarget.com/tip/6-Amazon-SageMaker-capabilities-developers-should-know-about>
- <https://aws.amazon.com/de/machine-learning/>
- <https://datasolut.com/crisp-dm-standard/#Business-Understanding>
- https://kedro.readthedocs.io/en/stable/01_introduction/01_introduction.html
- https://kedro.readthedocs.io/en/stable/03_tutorial/06_visualise_pipeline.html
- <https://docs.bentoml.org/en/latest/quickstart.html#run-on-google-colab>
- <https://www.fiddler.ai/ml-monitoring>

